


Deep Neural Networks for Doubly Robust Estimation with Nonprobability Survey Samples

Based on Dai et al. 2026

Motivation: integrating two imperfect data sources

- **Nonprobability sample** S_A : rich outcome y , but unknown selection mechanism.
- **Probability sample** S_B : design weights and representative covariates x , but often no y .
- Goal: estimate the finite-population mean

$$\mu_y = \frac{1}{N} \sum_{i=1}^N y_i.$$



Formal setup and notation

- Finite population $U = \{1, \dots, N\}$; target parameter: $\mu_y = \frac{1}{N} \sum_{i=1}^N y_i$.
- **Nonprobability sample** S_A , size n_A : observe (x_i, y_i) . Selection indicator $R_i = I(i \in S_A)$, with $\pi_i^A = P_q(R_i = 1 \mid x_i, y_i)$.
- **Probability sample** S_B , size n_B : observe (x_i, d_i^B) , where $d_i^B = 1/\pi_i^B$ and $\pi_i^B = P(i \in S_B)$ are the first-order inclusion probabilities.
- **Ignorability assumption**: $\pi_i^A = P(R_i = 1 \mid x_i)$ — selection depends only on observed x , not on y .

Three classical strategies

- 1 Model the outcome: $m(x_i) = E(y_i \mid x_i) \rightarrow$ predict y for units in S_B .
- 2 Model the selection: $\pi^A(x) = P(R = 1 \mid x) \rightarrow$ inversely weight units in S_A .
- 3 Combine both models (doubly robust).

Method 1: Regression / mass imputation (REG)

Procedure

- 1 Fit outcome model $m(x_i) = E(y_i | x_i)$ using S_A (e.g., OLS).
- 2 Predict $\hat{y}_i = m(x_i, \hat{\beta})$ for every unit in S_B .
- 3 Estimate the population mean by

$$\hat{\mu}_{REG} = \frac{1}{\hat{N}_B} \sum_{i \in S_B} d_i^B \hat{y}_i, \quad \hat{N}_B = \sum_{i \in S_B} d_i^B.$$

Critical assumption

$m(x)$ is correctly specified. If $E(y | x)$ contains interactions, thresholds, or nonlinearities and a linear model is used, $\hat{\mu}_{REG}$ is inconsistent — no amount of design weighting can salvage a misspecified outcome model.

Method 2: Inverse probability weighting (IPW)

Procedure

- 1 Pool S_A and S_B ; define $R_i = I(i \in S_A)$.
- 2 Fit propensity model $\pi^A(x) = P(R = 1 | x)$, typically logistic: $\text{logit}[\pi^A(x)] = \alpha_0 + x^\top \alpha$.
- 3 Weight each unit in S_A inversely to its selection probability:

$$\hat{\mu}_{IPW} = \frac{1}{\hat{N}_A} \sum_{i \in S_A} \frac{1}{\hat{\pi}_i^A} y_i, \quad \hat{N}_A = \sum_{i \in S_A} \frac{1}{\hat{\pi}_i^A}.$$

Critical assumption

$\pi^A(x)$ follows the assumed parametric form (e.g., linear logit). If the true selection mechanism involves **nonlinearities, interactions, or unknown functional forms**, the logistic model is misspecified and IPW is inconsistent.

Method 3: Doubly robust (DR) estimation

Formula

$$\hat{\mu}_{DR} = \frac{1}{\hat{N}_A} \sum_{i \in S_A} \frac{1}{\hat{\pi}_i^A} \{y_i - m(x_i, \hat{\beta})\} + \frac{1}{\hat{N}_B} \sum_{i \in S_B} \frac{1}{\pi_i^B} m(x_i, \hat{\beta}).$$

Double robustness property

$\hat{\mu}_{DR}$ is consistent if **either** the outcome model $m(x)$ **or** the propensity model $\pi^A(x)$ is correctly specified — not necessarily both. This is a genuine improvement over REG and IPW alone.

Remaining vulnerability

When **both** working models are misspecified — e.g., the outcome regression is linear but $E(y | x)$ is nonlinear, *and* the propensity is logistic but the true $\pi^A(x)$ is nonlinear — the DR estimator can also be inconsistent.

Parametric double robustness $\neq$ nonparametric protection.

Key idea of the paper

The paper replaces the parametric propensity-score model by a DNN model:

$$\log \frac{\pi_i^A}{1 - \pi_i^A} = g_0(x_i),$$

where g_0 is an unknown nonparametric function.



The DNN is trained using information from both the nonprobability sample and the reference probability sample.

Why deep neural networks?

- **Universal approximation:** multilayer feedforward networks can approximate broad classes of continuous functions on compact sets.
- **Curse of dimensionality:** under smoothness and structural (e.g. composite Hölder) assumptions, sparse DNNs can achieve the optimal minimax rate without suffering from the curse of dimensionality.
- **Scalability:** compared with classical nonparametric methods (e.g. kernel smoothing), DNNs scale better in moderate-to-high dimensions when the target function has suitable structure.
- **Nonlinear selection:** participation mechanisms in nonprobability samples often involve unknown interactions and nonlinearities — precisely the setting where DNNs excel.

Key insight

The parametric logit model $\text{logit}[\pi^A(x)] = \alpha_0 + x^\top \alpha$ is a **special case** of the DNN model. The DNN relaxes the linear-in-parameters restriction and lets the data determine the functional form of the selection mechanism.

Pseudo-likelihood: derivation

Step 1: Population log-likelihood (infeasible)

If x_i were observed for all $i \in U$, the log-likelihood for the participation indicators R_i would be

$$\ell(g) = \sum_{i=1}^N [R_i \log \pi_i^A + (1 - R_i) \log(1 - \pi_i^A)] = \sum_{i \in S_A} \log \frac{\pi_i^A}{1 - \pi_i^A} + \sum_{i=1}^N \log(1 - \pi_i^A).$$

Step 2: Replace population total by HT estimator

Since x_i is unavailable outside $S_A \cup S_B$, the population sum $\sum_{i=1}^N \log(1 - \pi_i^A)$ cannot be computed. The **Horvitz–Thompson estimator** uses the probability sample S_B :

$$\sum_{i=1}^N \log(1 - \pi_i^A) \approx \sum_{i \in S_B} d_i^B \log(1 - \pi_i^A), \quad d_i^B = \frac{1}{\pi_i^B}.$$

Pseudo-likelihood: final form

Step 3: Nonparametric logit model

Under the DNN model $\log \frac{\pi_i^A}{1-\pi_i^A} = g(x_i)$, we have $\log(1 - \pi_i^A) = -\log\{1 + \exp[g(x_i)]\}$.

The pseudo-log-likelihood simplifies to

$$\ell^*(g) = \sum_{i \in S_A} g(x_i) - \sum_{i \in S_B} d_i^B \log\{1 + \exp[g(x_i)]\}$$

- First term: uses units from the nonprobability sample S_A .
- Second term: uses the probability sample S_B as a design-weighted proxy for the population.
- A sparse ReLU DNN approximates g_0 ; ADAM minimizes $-\ell^*(g)$.

DNN architecture and training

Network structure

A $(K + 1)$ -layer feedforward DNN with width vector $\mathbf{p} = (p_0, \dots, p_K, p_{K+1})$:

- $p_0 = r$ (input: covariate dimension); $p_{K+1} = 1$ (output: scalar logit).
- K hidden layers with ReLU activation: $\sigma(u) = \max\{0, u\}$.
- Sparse network class $\mathcal{G}(K, s, \mathbf{p}, D)$: total number of non-zero parameters $\leq s$, sup-norm $\leq D$.

Training

- Loss function: $-\ell^*(g)$ (negative pseudo-log-likelihood).
- Optimizer: ADAM (adaptive moment estimation).
- After training: $\hat{\pi}_i^A = \{1 + \exp[-\hat{g}(x_i)]\}^{-1}$.

Proposed estimators

After training the DNN,

$$\hat{\pi}_i^A = \{1 + \exp[-\hat{g}(x_i)]\}^{-1}.$$

DNN-assisted IPW

$$\hat{\mu}_{DIPW} = \frac{1}{\hat{N}_A} \sum_{i \in S_A} \frac{1}{\hat{\pi}_i^A} y_i, \quad \hat{N}_A = \sum_{i \in S_A} \frac{1}{\hat{\pi}_i^A}.$$

Deep doubly robust estimator

$$\hat{\mu}_{DDR} = \frac{1}{\hat{N}_A} \sum_{i \in S_A} \frac{1}{\hat{\pi}_i^A} \{y_i - m(x_i, \hat{\beta})\} + \frac{1}{\hat{N}_B} \sum_{i \in S_B} \frac{1}{\pi_i^B} m(x_i, \hat{\beta}).$$

Assumptions: identification and design

A1: Ignorability (Missing At Random)

$\pi_i^A = P(R_i = 1 \mid x_i, y_i) = P(R_i = 1 \mid x_i)$. Selection into S_A depends only on x_i , not on y_i . Hence $E(y \mid x) = E(y \mid x, R = 1)$, so the outcome model is estimable from S_A .

A2: Positivity

$\pi_i^A > 0$ for all $i \in U$. Every unit has a nonzero chance of entering S_A . Together with A1, this is the **strong ignorability** condition (Rosenbaum & Rubin, 1983).

A3: Conditional independence

$R_i \perp R_j \mid x_i, x_j$ for $i \neq j$. Membership indicators are independent across units given covariates — rules out cluster or network dependence in the selection process.

Assumptions: DNN complexity and regularity

B1: DNN architecture

The network depth K , width p_k , and sparsity s scale as

$$K = O(\log n), \quad s = O(n\gamma_n^2 \log n), \quad \min_k p_k \asymp \max_k p_k \asymp n.$$

Deep enough to approximate g_0 , sparse enough to avoid overfitting.

B2: Smoothness of g_0

$g_0 \in \mathcal{H}(q, \beta, \mathbf{d}, \tilde{\mathbf{d}}, M)$: a composite Hölder class. Each component has smoothness β_i and intrinsic dimension $\tilde{d}_i \leq d_i$. Sparse ReLU networks approximate such g_0 at the rate γ_n .

Assumptions: regularity conditions

Design conditions (C1, C4, C5)

- **C1:** $n_A/N \rightarrow f_A \in (0, 1)$, $n_B/N \rightarrow f_B \in (0, 1)$ — samples are of comparable order.
- **C4:** HT estimators are $\sqrt{n_B}$ -consistent for population totals of x_i , y_i , $m(x_i, \beta)$, $m'(x_i, \beta)$.
- **C5:** $N\pi_i^A/n_A$ and $N\pi_i^B/n_B$ bounded away from 0 and ∞ — like simple random sampling.

Technical conditions (C2–C3, C6–C9)

- **C2–C3:** $m(x, \beta)$ twice continuously differentiable in β with bounded derivatives.
- **C6:** Finite moments of y_i^2 , $\|x_i\|^3$ bounded; propensity-score information matrix positive definite.
- **C7:** Logistic link $\sigma(z) = (1 + e^{-z})^{-1}$ has globally bounded derivatives.
- **C8:** Uniform design-consistency of HT estimators over the sparse DNN class.
- **C9:** Estimated scores truncated to $[\kappa_n, 1 - \kappa_n]$ for numerical stability.

Theorem 1: Convergence rate of the DNN estimator

Setup

g_0 belongs to a composite Hölder class $\mathcal{H}(q, \beta, \mathbf{d}, \tilde{\mathbf{d}}, M)$ where each component has smoothness β_i and intrinsic dimension $\tilde{d}_i \leq d_i$. The overall rate is $\gamma_n = \max_{i=0, \dots, q} n^{-\tilde{\beta}_i / (2\tilde{\beta}_i + \tilde{d}_i)}$ with $\tilde{\beta}_i = \beta_i \prod_{k=i+1}^q (\beta_k \wedge 1)$.

Theorem 1

Under A1–A3, B1–B2, C1, C8: $\|\hat{g} - g_0\|_{L^2([0,1]^r)} = O_p(\gamma_n \log^2 n)$.

Key insight

The rate depends on the **intrinsic dimensions** $\tilde{d}_i$, not the ambient r . When $\tilde{d}_i \ll d_i$, DNNs circumvent the curse of dimensionality.

Theorem 2: Convergence of DIPW and DDR

Theorem 2

Under A1–A3, B1–B2, C1–C9 (including correct specification of g_0):

$$|\hat{\mu}_{DIPW} - \mu_y| = O_p(\gamma_n \log^2 n), \quad |\hat{\mu}_{DDR} - \mu_y| = O_p(\gamma_n \log^2 n).$$

Interpretation

Both estimators inherit the DNN convergence rate, which depends on the composite smoothness and intrinsic dimension of g_0 , not the raw covariate count.

Simulation design and main findings

- Population size: $N = 20000$; sample sizes: $n_A = 500$, $n_B = 1000$.
- True propensity score is nonlinear: interactions, squared terms, sine and log terms.
- Compared estimators: naive mean, REG, IPW, DR, DIPW, DDR.

Main patterns

- **TF**: outcome model correct, parametric propensity wrong. DIPW improves over IPW; DDR has small bias and MSE.
- **FF**: both parametric working models wrong. REG, IPW and DR deteriorate; DIPW and DDR remain more stable.

Important nuance

The strong FF performance is finite-sample evidence in this simulation, not a guarantee under arbitrary misspecification.

Empirical application

- Nonprobability data: 2015 Pew Research Center sample, $n_A = 9301$.
- Probability reference: 2015 BRFSS, $n_B = 441456$.
- Common covariates: age, gender, race, region, marital status, employment, education.
- Outcomes include neighbor trust, political participation, local voting, and drinking days.

Empirical interpretation

Adjusted estimators differ from the raw nonprobability mean, showing that the probability sample changes the target-population adjustment. Since the true population means are unknown, the application mainly illustrates implementation rather than definitive accuracy.

References I

-  Chen, Yilin, Peng Li, and Changbao Wu (2020). “Doubly Robust Inference with Nonprobability Survey Samples”. In: *Journal of the American Statistical Association* 115, pp. 2011–2021.
-  Dai, Yufang et al. (2026). *Deep Neural Networks for Doubly Robust Estimation with Nonprobability Survey Samples*. arXiv: 2605.28762 [math.ST]. URL: <https://arxiv.org/abs/2605.28762>.
-  Goodfellow, Ian, Yoshua Bengio, and Aaron Courville (2016). *Deep Learning*. Cambridge, MA: MIT Press.
-  Kingma, Diederik P. and Jimmy Ba (2014). “Adam: A Method for Stochastic Optimization”. In: *arXiv preprint arXiv:1412.6980*.
-  Leshno, Moshe et al. (1993). “Multilayer Feedforward Networks with a Nonpolynomial Activation Function Can Approximate Any Function”. In: *Neural Networks* 6, pp. 861–867.

References II

-  Nair, Vinod and Geoffrey E. Hinton (2010). “Rectified Linear Units Improve Restricted Boltzmann Machines”. In: *Proceedings of the 27th International Conference on Machine Learning (ICML-10)*, pp. 807–814.
-  Rosenbaum, Paul R. and Donald B. Rubin (1983). “The Central Role of the Propensity Score in Observational Studies for Causal Effects”. In: *Biometrika* 70, pp. 41–55.
-  Schmidt-Hieber, Johannes (2020). “Nonparametric Regression Using Deep Neural Networks with ReLU Activation Function”. In: *The Annals of Statistics* 48, pp. 1916–1921.
-  Zhong, Qixian, Jonas Mueller, and Jane-Ling Wang (2022). “Deep Learning for the Partially Linear Cox Model”. In: *The Annals of Statistics* 50, pp. 1348–1375.